

# Ch 5.1.1-2: Leave One Out Cross-validation

## Lecture 9 - CMSE 381

Michigan State University

::

Dept of Computational Mathematics, Science & Engineering

February 7, 2024

# Covered in this lecture

- LOO CV
- Outliers
- Leverage statistic
- k-fold CV

## Section 1

Validation set

# What's the problem?

- How well is my ML method doing? *Model Assessment*
- Which method is best for our data?
- How many features should I use? Which ones? *Model selection*
- What is the uncertainty in the learned parameters?

*First, Cross Validation. Then Bootstrap.*

# Training Error vs Testing Error

## Training Error

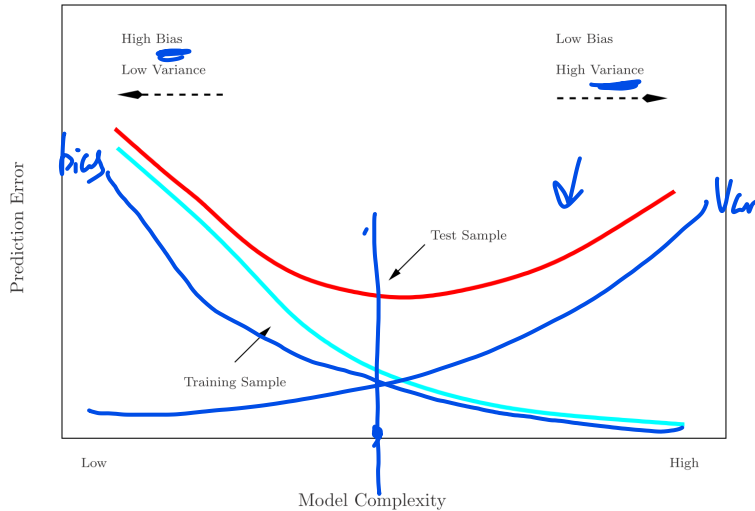
*Easily calculated by applying stats learning method to the observations used in its training*

- Massive under estimate
- overfitting!


## Testing Error

*Avg error that results from using a statistical learning method to predict the response on a new observation, one that was not used in training the method.*

# Model tradeoffs

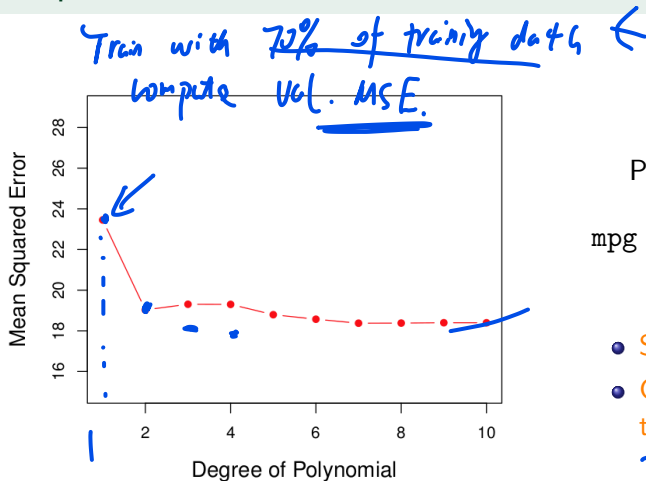


# Validation set approach



- Divide randomly into two parts:
  - ▶ Training set
  - ▶ Validation/Hold-out/Testing set
- Fit model on training set
- Use fitted model to predict response for observations in the test set
- Evaluate quality (e.g. MSE)
- This example has 30 data points
- Keep out 30% = 9 for testing in yellow

## Example with the auto data




Predicting mpg using horsepower:

$$\text{mpg} = \beta_0 + \beta_1 \text{hp} + \beta_2 \text{hp}^2 + \dots + \beta_p \text{hp}^p$$


- Shows the TEST error
- Can conclude that linear isn't enough to model the data



## Rinse and repeat



## Again example with auto data




- Repeat the procedure with different splits
- Highly variable results, no consensus about the error
- Further complicated: Machine learning methods do worse when trained on fewer observations, so validation set error rate tends to overestimate test error rate when fitted on the whole data set.

## Section 2

### Leave-One-Out Cross-Validation (LOOCV)

# The idea



# The idea in mathy words

- Remove  $(x_1, y_1)$  for testing.
- Train the model on  $n - 1$  points:  
→  $\{(x_2, y_2), \dots, (x_n, y_n)\}$
- Calculate  $\text{MSE}_1 = (y_1 - \hat{y}_1)^2$

Return the score:

$$CV_{(n)} = \frac{1}{n} \sum_{i=1}^n \text{MSE}_i$$

- Remove  $(x_2, y_2)$  for testing.
- Train the model on  $n - 1$  points:  
→  $\{(x_1, y_1), (x_3, y_3), \dots, (x_n, y_n)\}$
- Calculate  $\text{MSE}_2 = (y_2 - \hat{y}_2)^2$
- Rinse and repeat

# LOOCV Pros and Cons

## Advantages:


- Using almost all data every time
- Stable results: no randomness so always the same answer on a fixed data set
- Evaluation is performed with more training data

## Disadvantages:

- Computationally expensive; fit the model  $n$  times
- Doesn't shake up the data enough: estimates from each fold are highly correlated so their averages have high variance.

*Next goal: Speed up the LOOCV, but to do that, we need to get some terminology first*

## Again example with auto data



- There is no randomness. Same every time.


## Section 3

The one time you can cheat (by not computing every model fit)






# Outliers

An *outlier* is a point for which  $y_i$  is far from the value predicted by the model.  $|y_i - \hat{y}_i|$  is large  
*From Ch 3.3, stuff skipped earlier*



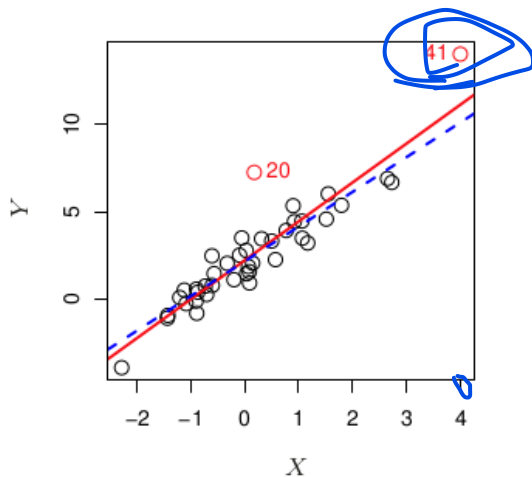
- Red line: Regression with all data
- Blue dashed: Regression after removing outlier 20
- Even though this is an outlier, didn't have much effect in this case.

# Residuals



Residuals are  $|y_i - \hat{y}_i|$ . Problem with no clear way to decide which is far enough to constitute outlier. Studentized version: divide the residual by its std error. Common choice is  $\geq 3$  is outlier.



# High Leverage



Observations with *high leverage* have an unusual value for  $x_i$ .

- Same data as previous but with extra data point 41
- Easy to find in  $p = 1$  case, as at left, since just look for big  $x$  value.
- Red line: fit all data
- Blue dashed line: fit after removing high leverage point 41
- Point out big movement in linear regression line

# Leverage statistic



Version for  $p = 1$

$$h_i = \frac{1}{n} + \frac{(x_i - \bar{x})^2}{\sum_{j=1}^n (x_j - \bar{x})^2}$$

*Generalizes to higher  $p$ , but not defined here. Compy can do it.*

# Leverage statistic properties

$$h_i = \frac{1}{n} + \frac{(x_i - \bar{x})^2}{\sum_{j=1}^n (x_j - \bar{x})^2}$$

- Increases with distance of  $x_i$  from  $\bar{x}$ .
- Always between  $1/n$  and 1
- Average leverage for all observations is always  $(p+1)/n$  (previous page this is  $2/41 = 0.049$ )
- If leverage stat is way bigger than  $(p+1)/n$ , then likely high leverage.
- 41 on previous page is both high leverage and outlier, so bad all around
- 20 on previous page had little effect on linear regression because low leverage even though it's an outlier

# Speeding up LOOCV

**Warning:** This only works for least squares linear or polynomial regression.


*fit the model on  $\{(x_i, y_i)\}_{i=1}^n$*   
 *$\hat{y}_i = \hat{f}(x_i)$*

$$\frac{1}{n} \sum_{i=1}^n \text{MSE}_i = CV_{(n)} = \frac{1}{n} \sum_{i=1}^n \left( \frac{y_i - \hat{y}_i}{1 - h_i} \right)^2$$


- Left side is original LOOCV computation
- Right side uses  $\hat{y}_i$  fit on the whole data set

- Yay, only one model fitting!
- Looks similar to regular MSE

Validation set



LOO-CV



## LOO-CV Score

$$CV_{(n)} = \frac{1}{n} \sum_{i=1}^n \text{MSE}_i$$

## Cheap trick for regression

$$CV_{(n)} = \frac{1}{n} \sum_{i=1}^n \left( \frac{y_i - \hat{y}_i}{1 - h_i} \right)^2$$

## Section 4

### $k$ -fold CV

•  $\eta_1, \dots, \eta_n$  are i.i.d  $N(0,1)$   $\leftarrow$







$$\rightarrow \text{Var} \left( \frac{\eta_1 + \dots + \eta_n}{n} \right) = \frac{\text{Var}(\eta_i)}{n} = \frac{1}{n} \leftarrow$$

•  $\eta_1 = \eta_2 = \dots = \eta_n$  ,  $\eta_i \sim N(0,1)$

$$\rightarrow \text{Var} \left( \frac{\eta_1 + \dots + \eta_n}{n} \right) < \text{Var}(\eta_i) = 1 \leftarrow$$




# Approximations of Test Error



# Definition of $k$ -fold CV



- Randomly split data into  $k$ -groups (folds)
- Approximately equal sized For the sake of notation, say each set has  $\ell$  points
- Remove  $i$ th fold  $U_i$  and reserve for testing.
- Train the model on remaining points
- Calculate
$$\text{MSE}_i = \frac{1}{\ell} \sum_{(x_j, y_j) \in U_i} (y_j - \hat{y}_j)^2$$
- Rinse and repeat



Return



$$CV_{(k)} = \frac{1}{k} \sum_{i=1}^k \text{MSE}_i$$

## Comparison with simulated data: Ex 3





- Blue: true test MSE is shown in blue
- Black dashed: the LOOCV estimate
- Orange: 10-fold CV
- The crosses indicate the minimum of each of the MSE curves.
- Nearly the same

# Comparison with simulated data: Ex 1



- Blue: true test MSE is shown in blue
- Black dashed: the LOOCV estimate
- Orange: 10-fold CV
- The crosses indicate the minimum of each of the MSE curves.
- **Underestimates the true test MSE**

## Comparison with simulated data: Ex 2



*Close for lower flexibility,  
then overestimates at  
higher ranges*

# Takeaways from the examples

- In all plots, LOOCV and 10-fold CV are similar
- Sometimes really only care about the minimum point in the test MSE because want to get the right degree of freedom for a model. All examples have approximately same x-coord for that

# Bias-Variance Tradeoff: Bias

$$\underline{E(y_0 - \hat{f}(x_0))^2} = \underline{\text{Var}(\hat{f}(x_0))} + \underbrace{\left[ \text{Bias}(\hat{f}(x_0)) \right]^2}_{\text{Bias}} + \underline{\text{Var}(\varepsilon)}$$

- 1.... Validation set overestimates test error b/c used small subset of data
- 3....  $k$ -fold gives medium level of bias b/c training set has approximately  $(k-1)n/k$  observations
- 2.... LOOCV gives approximately unbiased estimate since uses almost all data every time
- soooooooooo LOOCV better than  $k$ -fold just for bias

# Bias-Variance Tradeoff: Variance

$$E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon)$$

- LOOCV higher variance:
  - ▶ avg the outputs of  $n$  fitted models, but all on almost identical observations
  - ▶ high correlation with each other
- $k$ -fold
  - ▶  $k$  fitted models somewhat less correlated with each other
  - ▶ mean of many highly correlated quantities has higher variance than those not correlated
- Empirically,  $k = 5$  or  $10$  has been shown to be a happy medium